

This Is How Stripe Does Rate Limiting to Build Scalable APIs

#9: Read Now - Awesome Rate Limiter (4 minutes)



NEO KIM
SEP 28, 2023



37



3



2

Get the powerful template to approach system design for FREE on newsletter sign up

Subscribe

This post outlines how Stripe does rate limiting. If you want to learn more, scroll bottom and find the references.

- [Share this post](#) & I'll send you some rewards for the referrals.

Rate Limiter

A rate limiter is important to building a scalable API. Because it prevents bad use from abusing the API.

A rate limiter keeps a *counter* on the number of requests received. And reject a request if the threshold exceeds. Requests are rate-limited at the user or IP address level.

And it is a good choice if a change in the pace of the requests *doesn't* affect the experience.

Other potential *reasons* to rate limit are:


- Prevent low-priority traffic from affecting high-priority traffic
- Prevent service degradation

Rejecting low-priority requests under heavy load is called *load shedding*.

This post outlines how Stripe scales their API with the rate limiter. Consider [share this post](#) with someone who wants to study scalability patterns.

Rate Limiter Workflow

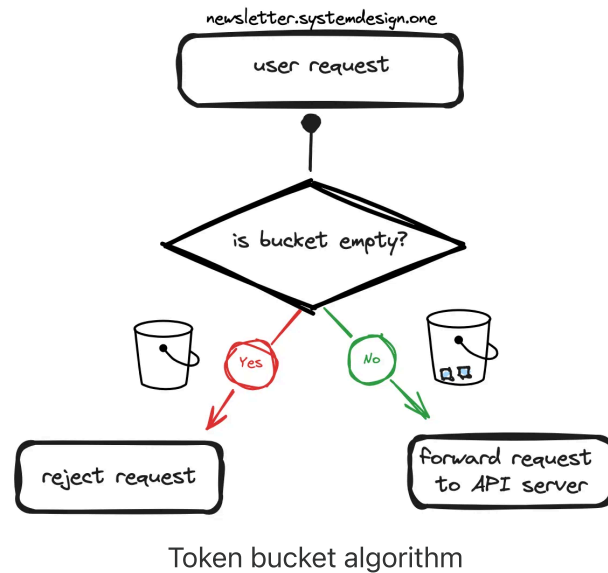
The rules for rate limiting are predefined. And here is the rate limiter workflow:



1. Check rate limiter rule
2. Reject a request if the threshold exceeds
3. Otherwise let the request pass-through

Rate Limiter Implementation

Stripe uses the [token bucket algorithm](#) to do rate limiting. Here is a quick overview of this algorithm:



Imagine there is a bucket filled with tokens. Every request must pick a token from bucket to pass through.

No requests get a token if the bucket is empty. So, further requests get rejected.

And tokens get refilled at a steady pace.

Other popular rate-limiting algorithms are sliding windows and leaky buckets.

They used Redis to build the rate limiter. Because it is in-memory and provides low latency.

Things they considered when implementing the rate limiter are:

- Quality check rate limiter logic and allow bypass on failures
- Show a clear response to the user: status code 429 - too many requests or service unavailable
- Enable panic mode on the rate limiter. This allowed switching it off on failure
- Set up alerts and monitoring
- Tune rate limiter to match traffic patterns

It's difficult to *rate-limit a distributed system*. Because each request from a single user might not hit the same server. I don't know how Stripe solved this problem. But I

a potential solution.

Redirect the traffic from an IP address to the same data center using DNS. And c
an isolated rate limiter in each data center.

Yet a new TCP connection might hit a different server within a data center.

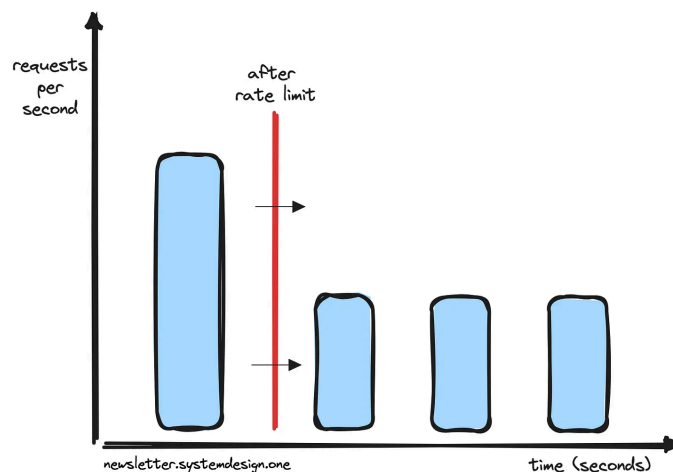
So set up a caching proxy ([twemproxy](#)) in each data center. Because it allows to
state across many servers. Put another way, many servers share a single cache:
limiter.

And use [consistent hashing](#) to reduce key redistribution on changing load (cluster
resize).

Rate Limiting Types

Stripe categorizes rate limiting into 4 types:


1. Request Rate Limiter



Request Rate Limiter

Each user gets n requests per second. This rate limiter type acts as the first line
defense for an API. And it is the most popular type.


2. Concurrent Requests Rate Limiter



Concurrent Requests Rate Limiter

The number of concurrent requests that are in progress is rate-limited. This protects resource-intensive API. And **prevents resource contention**.


3. Fleet Usage Load Shedder



Fleet Usage Load Shedder

The critical APIs reserve 20% of computing capacity. And requests to *non-critical* get rejected if the critical API doesn't get 20% of the resources.

4. Worker Utilization Load Shedder



Worker Utilization Load Shedder

The non-critical traffic gets shed on server overload. And it gets re-enabled after delay. This rate limiter type is the last line of defense for an API.

👋 PS - Are you unhappy at your current job?

While preparing for system design interviews to get your dream job can be stressful.

Don't worry, I'm working on content to help you pass the system design interview make it easier - you spend only a few minutes each week to go from 0 to 1. Yet your subscription fees will be higher than current pledge fees.

So pledge now to get access at a lower price.


"An excellent newsletter to learn system design through practical case studies."

Franco

Consider subscribing to get simplified case studies delivered straight to your inbox

Type your email...

Subscribe



Follow me on [LinkedIn](#) | [YouTube](#) | [Threads](#) | [Twitter](#) | [Instagram](#) | [Bluesky](#)

Thank you for supporting this newsletter. Consider sharing this post with your friends and get rewards. Y'all are the best.

1 Referral



Design Gurus
25% Discount Coupon

2 Referrals




Interview Mistakes
to Avoid PDF

3 Referrals



Popular
Interview Questions PDF



Tech Stack Evolution at Levels.fyi

NK • 24 SEPTEMBER 2023

[Read full story](#) →



11 Reasons Why YouTube Was Able to Support 100 Million Video Views a Day With Only 9 Engineers

NK • 16 SEPTEMBER 2023

[Read full story](#) →

Are you interested in deepening your knowledge of TDD and DDD? Consider subscribing to [Crafting Tech Teams Newsletter](#) from [Denis Čahuk](#). He has a great archive of information in his newsletter.

Word-of-mouth referrals like yours help this community grow - *Thank you.*



Samuel Path @smlpth · 12m

I enjoy the System Design Newsletter (@systemdesign42). Every few days I get a short article explaining what approach some tech company used to overcome a real-world challenge.

It's small and regular. Ideal to slowly grow one's understanding over time.

[Feedback from an awesome reader](#)

Get featured in the newsletter: Write your feedback on this post. And tag me on [Twitter](#), [LinkedIn](#), and [Substack Notes](#). Or, you can reply to this email with any feedback.

References

- <https://stripe.com/blog/rate-limiters>
- <https://www.cloudflare.com/en-gb/learning/bots/what-is-rate-limiting/>
- <https://blog.cloudflare.com/counting-things-a-lot-of-different-things/>

Subscribe to The System Design Newsletter

By Neo Kim · Launched 2 years ago

Weekly newsletter to help you become good at work

[Subscribe](#)

By subscribing, I agree to Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).



37 Likes · 2 Restacks

[← Previous](#)[Next →](#)

Discussion about this post

[Comments](#)[Restacks](#)

Write a comment...



Anton Zaides · 28 Sept 2023

♥ Liked by Neo Kim

If you implement the simple request rate limiter - one important part is how to communicate with clients. I had 2 cases in the last year where we breached the limit:

1. A free 3rd party api. We had a limit of 500 requests per day (which was not written down) and we didn't pass for a long time. Then we suddenly started to get 403 errors, which immediately made us think our token was revoked or expired. Only after it started to work at 12:00, did we understand we breached the limit, and for some reason, they decided to throw unauthorized errors instead of the standard 429.
2. A paid api, that we heavily use. Recently we reached their limit (100 calls/minute), and started to get 429 responses. The good part, is that in the header we got the time when our limit will be reset. This allowed us to implement internal queueing of the request, without noisy errors. We had a queue of requests, and when the time came, we stored the requests, and then we continued at the allowed time until we hit it again. In rare peaks, this solution is perfect for us.

♥ LIKE (6) 💬 REPLY

2 replies by Neo Kim and others

2 more comments...

© 2025 Neo Kim · [Publisher Privacy](#)

[Substack](#) · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture